

Cross-model replication: the Mahalanobis result holds, the entropy result does not

Same 19 output statistics throughout, so the comparison is architecture-independent. The fuzzy rule beats full-covariance Mahalanobis on both models (+0.036 and +0.032, both 8/8 seeds, $p = 0.008$) but beats mean entropy only on SmolLM2 long-form — it ties on Qwen and loses on the short-factual task.

